

```
In [1]: import os
img_dir = '/tmp/nst'
if not os.path.exists(img_dir):
    os.makedirs(img_dir)
!wget --quiet -P /tmp/nst/ https://upload.wikimedia.org/wikipedia
/commons/d/d7/Green_Sea_Turtle_grazing_seagrass.jpg
!wget --quiet -P /tmp/nst/ https://upload.wikimedia.org/wikipedia
/commons/0/0a/The_Great_Wave_off_Kanagawa.jpg
!wget --quiet -P /tmp/nst/ https://upload.wikimedia.org/wikipedia
/commons/b/b4/Vassily_Kandinsky%2C_1913_-_Composition_7.jpg
!wget --quiet -P /tmp/nst/ https://upload.wikimedia.org/wikipedia
/commons/0/00/Tuebingen_Neckarfront.jpg
!wget --quiet -P /tmp/nst/ https://upload.wikimedia.org/wikipedia
/commons/6/68/Pillars_of_creation_2014_HST_WFC3-UVIS_full-res_den
oised.jpg
!wget --quiet -P /tmp/nst/ https://upload.wikimedia.org/wikipedia
/commons/thumb/e/ea/Van_Gogh_-_Starry_Night_-_Google_Art_Project.
jpg/1024px-Van_Gogh_-_Starry_Night_-_Google_Art_Project.jpg
```

```
In [2]: import matplotlib.pyplot as plt
import matplotlib as mpl
mpl.rcParams['figure.figsize'] = (10,10)
mpl.rcParams['axes.grid'] = False

import numpy as np
from PIL import Image
import time
import functools
```

```
In [3]: %tensorflow_version 1.x
import tensorflow as tf

from tensorflow.python.keras.preprocessing import image as kp_ima
ge
from tensorflow.python.keras import models
from tensorflow.python.keras import losses
from tensorflow.python.keras import layers
from tensorflow.python.keras import backend as K
```

TensorFlow 1.x selected.

```
In [4]: tf.enable_eager_execution()
print("Eager execution: {}".format(tf.executing_eagerly()))
```

Eager execution: True

```
In [5]: content_path = '/tmp/nst/Green_Sea_Turtle_grazing_seagrass.jpg'
style_path = '/tmp/nst/The_Great_Wave_off_Kanagawa.jpg'
```

```
In [6]: def load_img(path_to_img):
    max_dim = 512
    img = Image.open(path_to_img)
    long = max(img.size)
    scale = max_dim/long
    img = img.resize((round(img.size[0]*scale), round(img.size[1]*scale)), Image.ANTIALIAS)

    img = kp_image.img_to_array(img)

    img = np.expand_dims(img, axis=0)
    return img
```

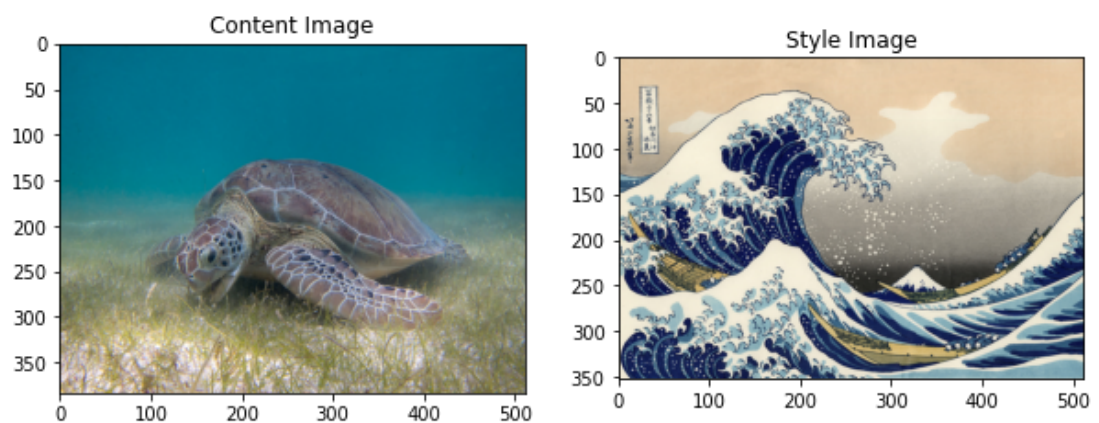
```
In [7]: def imshow(img, title=None):
    out = np.squeeze(img, axis=0)
    out = out.astype('uint8')
    plt.imshow(out)
    if title is not None:
        plt.title(title)
    plt.imshow(out)
```

```
In [8]: plt.figure(figsize=(10,10))

content = load_img(content_path).astype('uint8')
style = load_img(style_path).astype('uint8')

plt.subplot(1, 2, 1)
imshow(content, 'Content Image')

plt.subplot(1, 2, 2)
imshow(style, 'Style Image')
plt.show()
```



```
In [9]: def load_and_process_img(path_to_img):
    img = load_img(path_to_img)
    img = tf.keras.applications.vgg19.preprocess_input(img)
    return img
```

```
In [10]: def deprocess_img(processed_img):
x = processed_img.copy()
if len(x.shape) == 4:
    x = np.squeeze(x, 0)
assert len(x.shape) == 3,
if len(x.shape) != 3:
    raise ValueError("Invalid input to deprocessing image")

x[:, :, 0] += 103.939
x[:, :, 1] += 116.779
x[:, :, 2] += 123.68
x = x[:, :, ::-1]

x = np.clip(x, 0, 255).astype('uint8')
return x
```

```
In [11]: content_layers = ['block5_conv2']

style_layers = ['block1_conv1',
                'block2_conv1',
                'block3_conv1',
                'block4_conv1',
                'block5_conv1'
                ]

num_content_layers = len(content_layers)
num_style_layers = len(style_layers)
```

```
In [12]: def get_model():
vgg = tf.keras.applications.vgg19.VGG19(include_top=False, weights='imagenet')
vgg.trainable = False

style_outputs = [vgg.get_layer(name).output for name in style_layers]
content_outputs = [vgg.get_layer(name).output for name in content_layers]
model_outputs = style_outputs + content_outputs

return models.Model(vgg.input, model_outputs)
```

```
In [13]: def get_content_loss(base_content, target):
return tf.reduce_mean(tf.square(base_content - target))
```

```

In [14]: def gram_matrix(input_tensor):
    channels = int(input_tensor.shape[-1])
    a = tf.reshape(input_tensor, [-1, channels])
    n = tf.shape(a)[0]
    gram = tf.matmul(a, a, transpose_a=True)
    return gram / tf.cast(n, tf.float32)

def get_style_loss(base_style, gram_target):
    height, width, channels = base_style.get_shape().as_list()
    gram_style = gram_matrix(base_style)

    return tf.reduce_mean(tf.square(gram_style - gram_target))# /
    (4. * (channels ** 2) * (width * height) ** 2)

In [16]: def get_feature_representations(model, content_path, style_path):
    content_image = load_and_process_img(content_path)
    style_image = load_and_process_img(style_path)
    style_outputs = model(style_image)
    content_outputs = model(content_image)

    style_features = [style_layer[0] for style_layer in style_outputs[:num_style_layers]]
    content_features = [content_layer[0] for content_layer in content_outputs[num_style_layers:]]
    return style_features, content_features

In [17]: def compute_loss(model, loss_weights, init_image, gram_style_features, content_features):
    style_weight, content_weight = loss_weights
    model_outputs = model(init_image)

    style_output_features = model_outputs[:num_style_layers]
    content_output_features = model_outputs[num_style_layers:]

    style_score = 0
    content_score = 0

    weight_per_style_layer = 1.0 / float(num_style_layers)
    for target_style, comb_style in zip(gram_style_features, style_output_features):
        style_score += weight_per_style_layer * get_style_loss(comb_style[0], target_style)

    weight_per_content_layer = 1.0 / float(num_content_layers)
    for target_content, comb_content in zip(content_features, content_output_features):
        content_score += weight_per_content_layer * get_content_loss(comb_content[0], target_content)

    style_score *= style_weight
    content_score *= content_weight

    loss = style_score + content_score
    return loss, style_score, content_score

```

```
In [18]: def compute_grads(cfg):  
    with tf.GradientTape() as tape:  
        all_loss = compute_loss(**cfg)  
        total_loss = all_loss[0]  
    return tape.gradient(total_loss, cfg['init_image']), all_loss
```

```

In [19]: import IPython.display

def run_style_transfer(content_path,
                      style_path,
                      num_iterations=1000,
                      content_weight=1e3,
                      style_weight=1e-2):

    model = get_model()
    for layer in model.layers:
        layer.trainable = False

    style_features, content_features = get_feature_representations
(model, content_path, style_path)
    gram_style_features = [gram_matrix(style_feature) for style_fea
ture in style_features]

    init_image = load_and_process_img(content_path)
    init_image = tf.Variable(init_image, dtype=tf.float32)
    opt = tf.train.AdamOptimizer(learning_rate=5, beta1=0.99, epsil
on=1e-1)

    iter_count = 1

    best_loss, best_img = float('inf'), None

    loss_weights = (style_weight, content_weight)
    cfg = {
        'model': model,
        'loss_weights': loss_weights,
        'init_image': init_image,
        'gram_style_features': gram_style_features,
        'content_features': content_features
    }

    num_rows = 2
    num_cols = 5
    display_interval = num_iterations/(num_rows*num_cols)
    start_time = time.time()
    global_start = time.time()

    norm_means = np.array([103.939, 116.779, 123.68])
    min_vals = -norm_means
    max_vals = 255 - norm_means

    imgs = []
    for i in range(num_iterations):
        grads, all_loss = compute_grads(cfg)
        loss, style_score, content_score = all_loss
        opt.apply_gradients([(grads, init_image)])
        clipped = tf.clip_by_value(init_image, min_vals, max_vals)
        init_image.assign(clipped)
        end_time = time.time()

        if loss < best_loss:
            best_loss = loss
            best_img = deprocess_img(init_image.numpy())

    if i % display_interval == 0:

```

```

start_time = time.time()
plot_img = init_image.numpy()
plot_img = deprocess_img(plot_img)
imgs.append(plot_img)
IPython.display.clear_output(wait=True)
IPython.display.display_png(Image.fromarray(plot_img))
print('Iteration: {}'.format(i))
print('Total loss: {:.4e}, '
      'style loss: {:.4e}, '
      'content loss: {:.4e}, '
      'time: {:.4f}s'.format(loss, style_score, content_score, time.time() - start_time))
print('Total time: {:.4f}s'.format(time.time() - global_start))
IPython.display.clear_output(wait=True)
plt.figure(figsize=(14,4))
for i,img in enumerate(imgs):
    plt.subplot(num_rows,num_cols,i+1)
    plt.imshow(img)
    plt.xticks([])
    plt.yticks([])


return best_img, best_loss

```

```

In [20]: best, best_loss = run_style_transfer(content_path,
                                             style_path, num_iterations=1000)


```





```
In [23]: Image.fromarray(best)
```

```
Out[23]:
```



```
In [24]: def show_results(best_img, content_path, style_path, show_large_f
         final=True):
         plt.figure(figsize=(10, 5))
         content = load_img(content_path)
         style = load_img(style_path)

         plt.subplot(1, 2, 1)
         imshow(content, 'Content Image')


         plt.subplot(1, 2, 2)
         imshow(style, 'Style Image')

         if show_large_final:
             plt.figure(figsize=(10, 10))


             plt.imshow(best_img)
             plt.title('Output Image')
             plt.show()
```



```
In [25]: show_results(best, content_path, style_path)
```




```
In [21]: best_starry_night, best_loss = run_style_transfer('/tmp/nst/Tuebi  
nngen_Neckarfront.jpg',  
                                                         '/tmp/nst/1024p  
x-Van_Gogh_-_Starry_Night_-_Google_Art_Project.jpg')
```





```
In [26]: show_results(best_starry_night, '/tmp/nst/Tuebingen_Neckarfront.jpg',  
                    '/tmp/nst/1024px-Van_Gogh_-_Starry_Night_-_Google_Art_Project.jpg')
```



```
In [ ]:
```